Agentic AI in mathematical research

A brief disclaimer

I am not an expert in this field. I have been exploring it out of curiosity and to stay up to date.

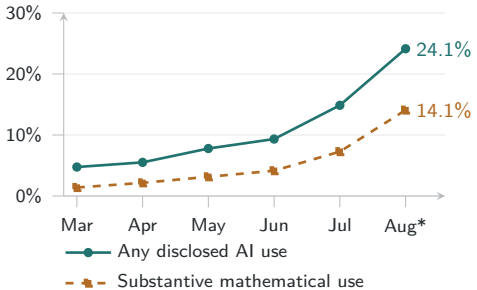
This talk neither argues for nor argues against using AI to write papers.

- A brief history and overview of mathematical research with LLMs, focusing on agentic approaches.
- What the literature tells us about successful approaches and their limitations.
- A case study where I build an AI research team to write a paper.

AI use and mathematical output

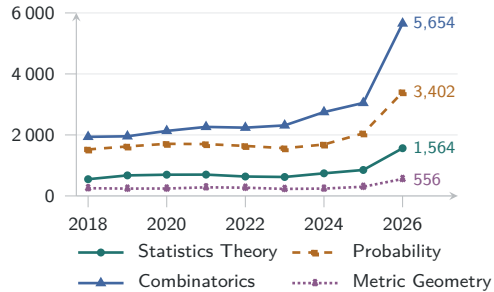
AI disclosures on arXiv

Share of math-related submissions, 2026



Submissions by field

New preprints, January–August each year

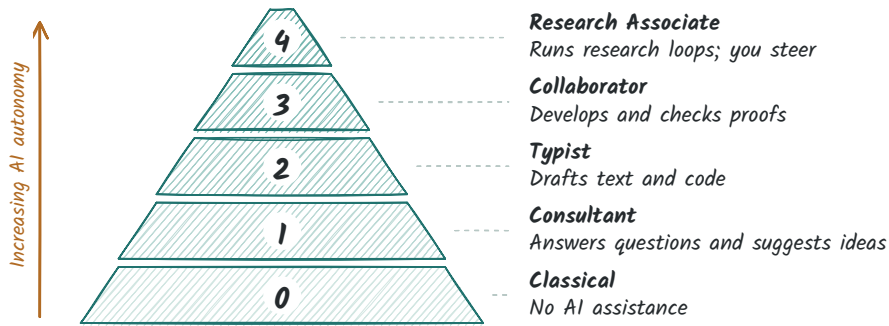


Left: 3,575 of 32,944 preprints disclosed AI use. *August covers 1–20.
Right: category counts for published arXiv papers.

Jin, Ke & Sui (2026), Table 1 Explore arXiv Math Pulse GitHub source.

(How) have you used AI in your research?

Which is the highest level you have used?



Adapted from Zimmer et al. (2026), *The Agentic Researcher*, Section 2.

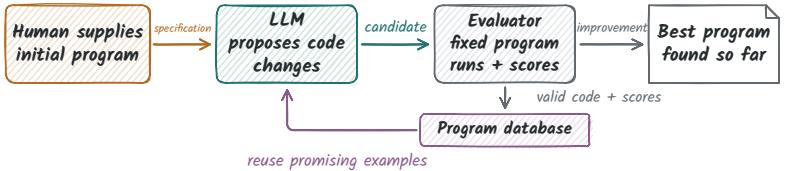
A short history of math research with LLMs

Agentic research systems over time

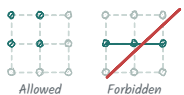
A brief history, with a focus on mathematics.

FunSearch Romera-Paredes et al., *Nature* (2024)

An evolutionary approach with fixed evaluation criteria implemented in computer programs.

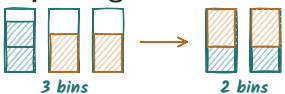


Larger cap sets



No three distinct points on a line in $\mathbb{F}_3^n$.
In eight dimensions: 496 to 512 points.

Online bin-packing heuristics



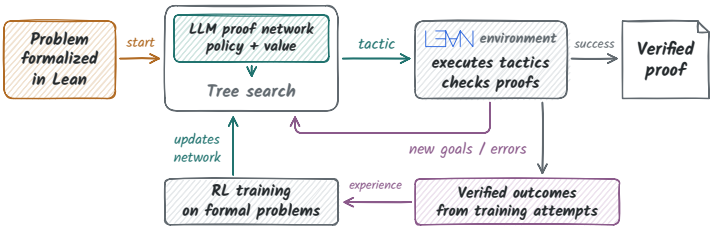
The paper showed that the model can improve a standard heuristic such as best fit.

AlphaEvolve (2025) later extended this approach: 4×4 complex matrix multiplication with 48 scalar multiplications, versus Strassen's 49.

Workflow: Fig. 1; results: Fig. 4 and Table 1. Romera-Paredes et al. (2024).

AlphaProof Hubert et al., *Nature* (2026)

Another way to verify and evaluate mathematics is through formal verification in Lean.



Function: the model reads Lean goals and hypotheses. As output, it proposes Lean tactics. Lean returns new proof states. This means it basically talks in Lean and gets Lean responses.

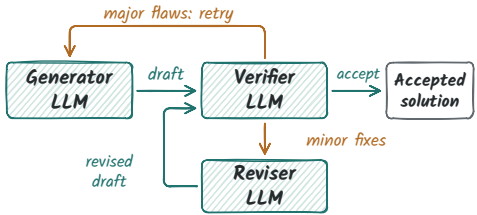
Additional learning: the model learns from proof attempts through reinforcement learning.

IMO 2024: 28/42 points, silver-medal level (AlphaProof: 3 problems; AlphaGeometry 2: 1).
Figs. 1–2; AlphaProof and Training sections. *Nature* 651, 607–613 (2026).

Aletheia Feng et al., arXiv (2026)

We can also use an independent LLM for verification. This does not guarantee correctness.

Gemini Deep Think agents



LLM judgment; no formal guarantee (Fig. 1).

Reported mathematical results

- Erdős: 2 full + 2 partial new solutions from 700 problems, checked by experts.

Other LLM research systems

The AI Scientist *Nature* (2026)
Ran ML experiments and wrote a paper that passed a workshop's first review round.

Google Co-Scientist *Nature* (2026)
Agents proposed and refined leukaemia drug combinations, validated in laboratory experiments.

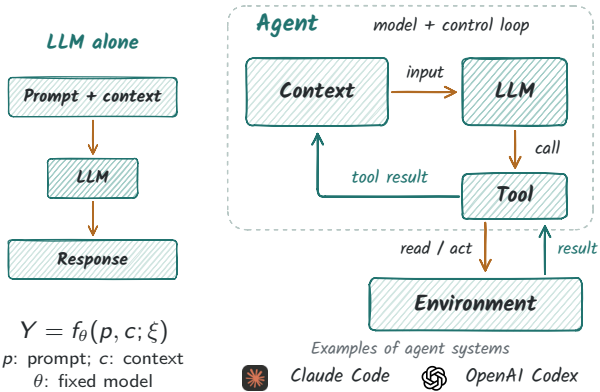
Virtual Lab *Nature* (2025)
LLM agents designed SARS-CoV-2 nanobodies with binding confirmed in laboratory tests.

OpenAI: Navier–Stokes report (2026)
Different approaches explored with 10,000 agents and human steering.

Agentic systems

From LLMs to research teams

Environment, context, tools and agents



$$Y = f_\theta(p, c; \xi)$$

p: prompt; c: context
θ: fixed model
ξ: randomness

Environment
The workspace an agent can access: files, data, code and external services.

Context
The information given to the LLM. This includes instructions, messages and tool results.

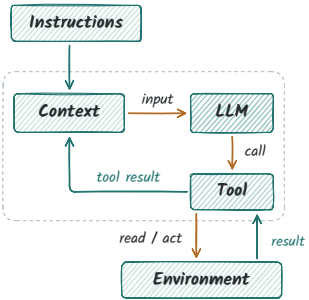
Tool
An interface to observe or change the environment, e.g. a PDF reader or Python.

Agent
An LLM-based system that uses tools and chooses its next step from results.

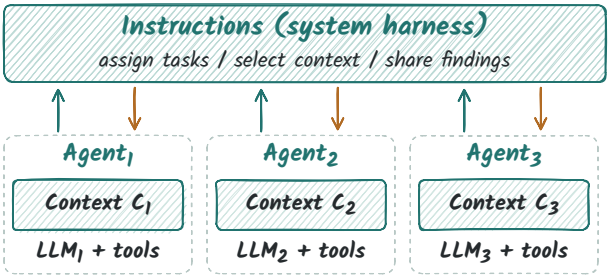
Based on OpenAI's Codex harness documentation and Claude Code's agentic loop.

Instructions and multi-agent systems

One agent



Several agents



Each agent has its own context. Instructions specify its task and how it should work.

Communication is managed through `instructions.md` or the system harness.

More agents are not automatically better.

Why manage context and communication?

An LLM produces text $Y = f_\theta(p, c; \xi)$ whose distribution depends strongly on the context.

1. Keep context focused Liu et al. (2024); Modarressi et al. (2025)
Give each agent the information it needs; long histories can obscure relevant details.
This helps reduce distraction from irrelevant context.

2. Give review a fresh context Sharma et al. (2023); Dhuliawala et al. (2023)
Review the draft and criteria without the writer's history; verify independently.
This helps avoid repeating errors from the draft.

3. Match the model to the task Ong et al., *RouteLLM* (2024)
Changing between stronger and cheaper models can improve the balance of cost and quality.

4. Share findings and avoid duplicate work Du et al. (2023); Wang et al. (2025)
Share useful results and redirect work that adds no new information.

5. Explore distinct approaches independently Zhu et al. (2026)
Especially important for mathematics since different starting points can lead to different proof strategies.

Agentic research: a worked example

What do we need?

- An agent system: in my case, Codex.
- A precise mathematical problem statement.
- Relevant related work: definitions, known results and useful methods.
- `instruction.md`: the task, constraints and expected output.
- `skills.md`: an index of reusable methods and checks.
- `AGENTS.md`: standing project instructions for the agents.

Step 1: Defining the problem

A **simple point process** X is a random set of distinct points in $\mathbb{R}^d$, finite in bounded regions.

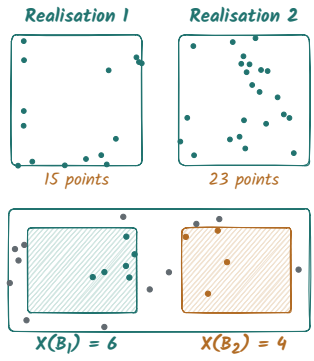
Its count in a bounded Borel region B is $X(B) = \#(X \cap B)$.

Stationarity means that shifting all points leaves the distribution unchanged:

$$X + t \overset{d}{=} X \quad \text{for every } t \in \mathbb{R}^d.$$

Homogeneous Poisson process, intensity $\lambda > 0$

- We demand that
$$X(B) \sim \text{Poisson}(\lambda|B|), \quad \mathbb{E}[X(B)] = \lambda|B|.$$
- Counts in disjoint regions are independent.
- Therefore, for $B_1 \cap B_2 = \emptyset$,
$$\mathbb{E}[X(B_1)X(B_2)] = \mathbb{E}[X(B_1)]\mathbb{E}[X(B_2)].$$



Step 1: Defining the problem

Under suitable assumptions, as disjoint sets B_1, B_2 shrink to x_1, x_2,

$$\frac{\mathbb{E}[X(B_1)X(B_2)]}{\mathbb{E}[X(B_1)]\,\mathbb{E}[X(B_2)]} \longrightarrow g_2(x_1, x_2).$$

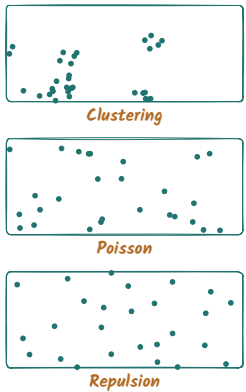
More generally, for disjoint sets $B_1, \ldots, B_n$ shrinking to $x_1, \ldots, x_n$,

$$\frac{\mathbb{E}[X(B_1)X(B_2)\cdots X(B_n)]}{\mathbb{E}[X(B_1)]\,\mathbb{E}[X(B_2)]\cdots\mathbb{E}[X(B_n)]} \longrightarrow g_n(x_1, x_2, \ldots, x_n).$$

This is the *n*-**point correlation function**: $g_n = \rho^{(n)}/\lambda^n$,
where $\rho^{(n)}$ is the density of expected ordered distinct tuples.

Compared with Poisson at the same intensity:

$g_n > 1$: more such tuples.
$g_n < 1$: fewer such tuples.
$g_n = 1$: the same local tuple density.



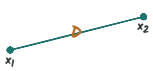
Step 1: The precise problem

The conjecture is motivated by observed spatial correlations in earthquake catalogues. See Kagan and Knopoff (1980); Kagan (1981a, 1981b).

For any intensity $\lambda_0 > 0$, there exists a stationary, simple, locally finite point process X on $\mathbb{R}^3$ with existing product densities, such that for some constants $C_2, C_3, C_4 > 0$,

$$g_2(x_1, x_2) = \frac{C_2}{D}, \qquad g_3(x_1, x_2, x_3) = \frac{C_3}{A}, \qquad g_4(x_1, x_2, x_3, x_4) = \frac{C_4}{V}$$

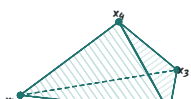
almost everywhere.



Distance D



Area A



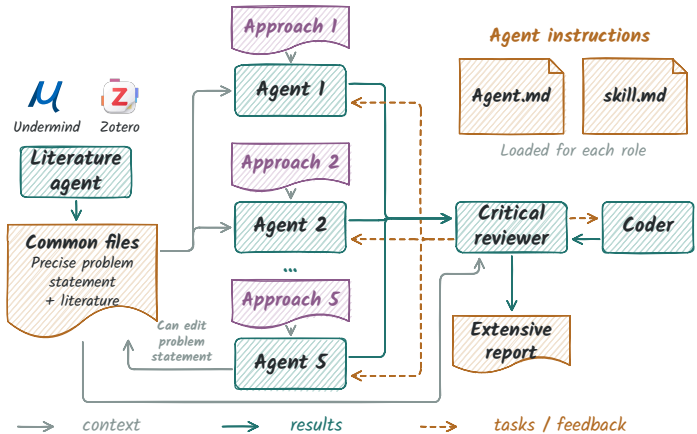
Volume V

Additional assumption: require these identities only for diameter $< M$

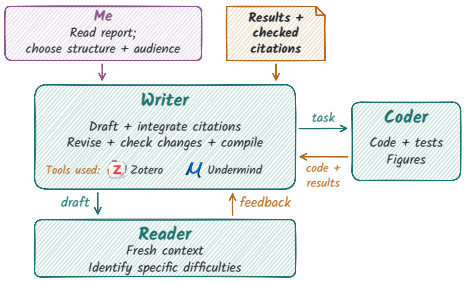
Step 2: Configuration/agent files and execution

Research approaches

1. Poisson cluster process
2. Cox process
3. Hawkes process
4. An approach of the agent's choice
5. Prove non-existence



Step 3: Writing the paper



- The paper's appendix also discusses the other approaches. They led to interesting ideas, even when they did not work.
- The writer gets access to Zotero and Undermind for citations and novelty checks.

Research phase: recorded usage

Work	Tokens (millions)	Est. USD
Proofer	69.74	$121.04
Reviewer	25.94	$54.47
Literature	3.40	$5.65
Total	**99.08**	**$181.15**

Writing phase: paper, simulations and review

Token category	Tokens	Est. USD
Uncached input	545,555	$5.46
Cached input	16,615,424	$16.62
Output (incl. reasoning)	116,258	$5.81
Total	**17,277,237**	**$27.88**

Standard GPT-6-Astra API pricing.

Results and personal feedback

Spheres and a cluster correction

Idea: I choose between two processes, each with probability 0.5. Therefore

$$g_n = \frac{\lambda_A^n g_{n,A} + \lambda_B^n g_{n,B}}{2\lambda_0^n}, \quad \lambda_0 = \frac{\lambda_A + \lambda_B}{2}.$$

Process A. Takes care of $g_4 \propto V^{-1}$ by placing non-interacting spheres with random radii. A homogeneous Poisson process on each sphere has the required g_4. We also know g_3 and g_2 here.

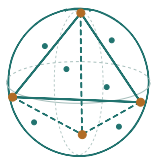
Process B. Ensures that $g_{4,B} = 0$ for diameters below M, using clusters of at most 3 points separated by more than M. $g_{2,B}$ and $g_{3,B}$ are designed according to the error in Process A via the cluster kernels

$$F_3 = q_3, \qquad F_2 = q_2 - \int q_3(\cdot, \cdot, z)\, dz.$$

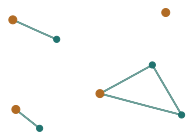
Here q_2, q_3 are the required product-density corrections.

Personal feedback

Aspect	Feedback
Mathematical correctness	Positive
Research and process documentation	Positive
Code writing	Positive
Text writing	Medium
Paper structuring	Negative
Mathematical explanation	Medium



Process A



Process B

Future of agentic research

What I expect to change

- Stronger models can increasingly plan, divide work into tasks, use tools and delegate to subagents.
- For a precise mathematical problem, one initial prompt may increasingly suffice. This can still mean substantial autonomous work, with no guarantee of success.
- I also expect explanations to become less sloppy and more useful for learning.

Where a custom harness still helps

- Keep the exact target, hypotheses and criteria for success fixed.
- Require separate review and formal verification.
- Keep human checkpoints for direction and interpretation.
- Record results, failed approaches, sources and dependencies.
- Set compute budgets, stopping rules and model choice.
- When writing papers with AI, set your preferences for writing and illustration styles.

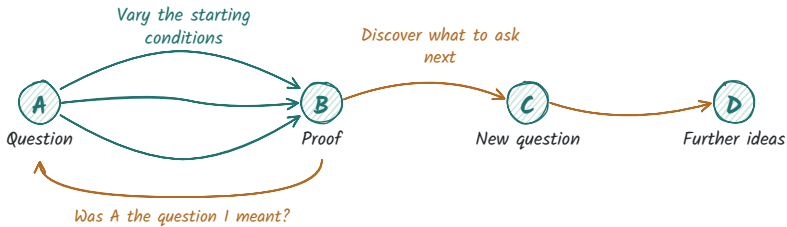
A reported example. Scientific American reports that almost all results in OpenAI's October release came from one prompt to one agent. Some may have needed retries; the full workflow has not been independently reproduced.
Scientific American, October 2026; OpenAI release details.

Personal take

What I have found useful in my research:

- Writing code.
- Managing my literature.
- Exploring ideas and proving theorems.
- Getting quick new insights.
- Visualising my science.
- Correcting my writing.

Personal take



Reaching B only helps me if I understand the mathematics.
Reaching B in different ways lets me think about where to go next, i.e. C.

My concern as a PhD student: outsourcing too much may bypass the learning.
A PhD is still an education. Producing results and developing understanding can come apart.

Related: *A Severe Misalignment of AI in Mathematics* (joint declaration, 11 September 2026).

Thank you!

Questions and discussion

Slides, experiments, documentation and the paper:
github.com/MathisRost/agentic-ai-math

